databricks

Delta Live Tables

Automatic reliable ETL on Delta Lake

Chris Hoshino-Fish

Lead Solutions Architect, Databricks since 2017

Specialize in Real-Time Data systems & Performance Engineering

Data Engineer since 2014

B.A. Computational Mathematics, UC Santa Cruz 2012

fish@databricks.com

UNIVERSITY OF CALIFORNIA SANTA CRUZ

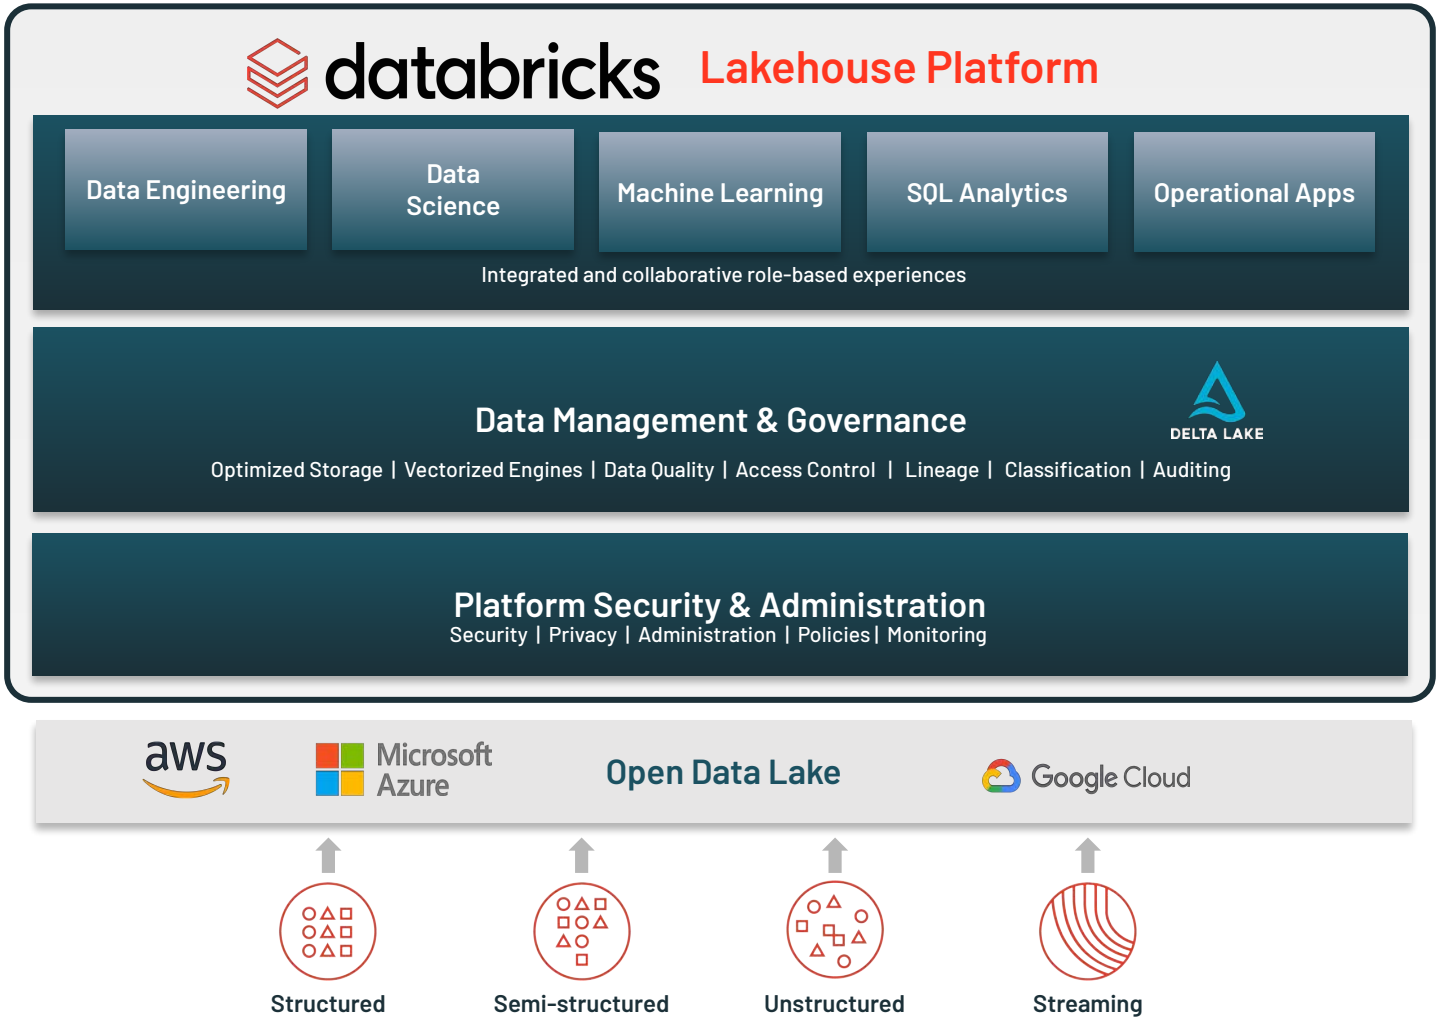
databricks
Lakehouse Platform
Data Engineering
Data Science
Machine Learning
SQL Analytics
Operational Apps
Integrated and collaborative role-based experiences
Data Management & Governance
DELTA LAKE
Optimized Storage | Vectorized Engines | Data Quality | Access Control | Lineage | Classification | Auditing
Platform Security & Administration
Security | Privacy | Administration | Policies | Monitoring
aws
Microsoft Azure
Open Data Lake
Google Cloud
Structured
Semi-structured
Unstructured
Streaming

Data Warehouses

Pros

- Great for Business Intelligence (BI) applications

Cons

- Limited support for Machine Learning (ML) workloads
- Proprietary systems with only a SQL interface

Data Lakes

Pros

- Supports ML
- Open formats and big ecosystem

Cons

- Poor support for BI
- Complex data quality problems

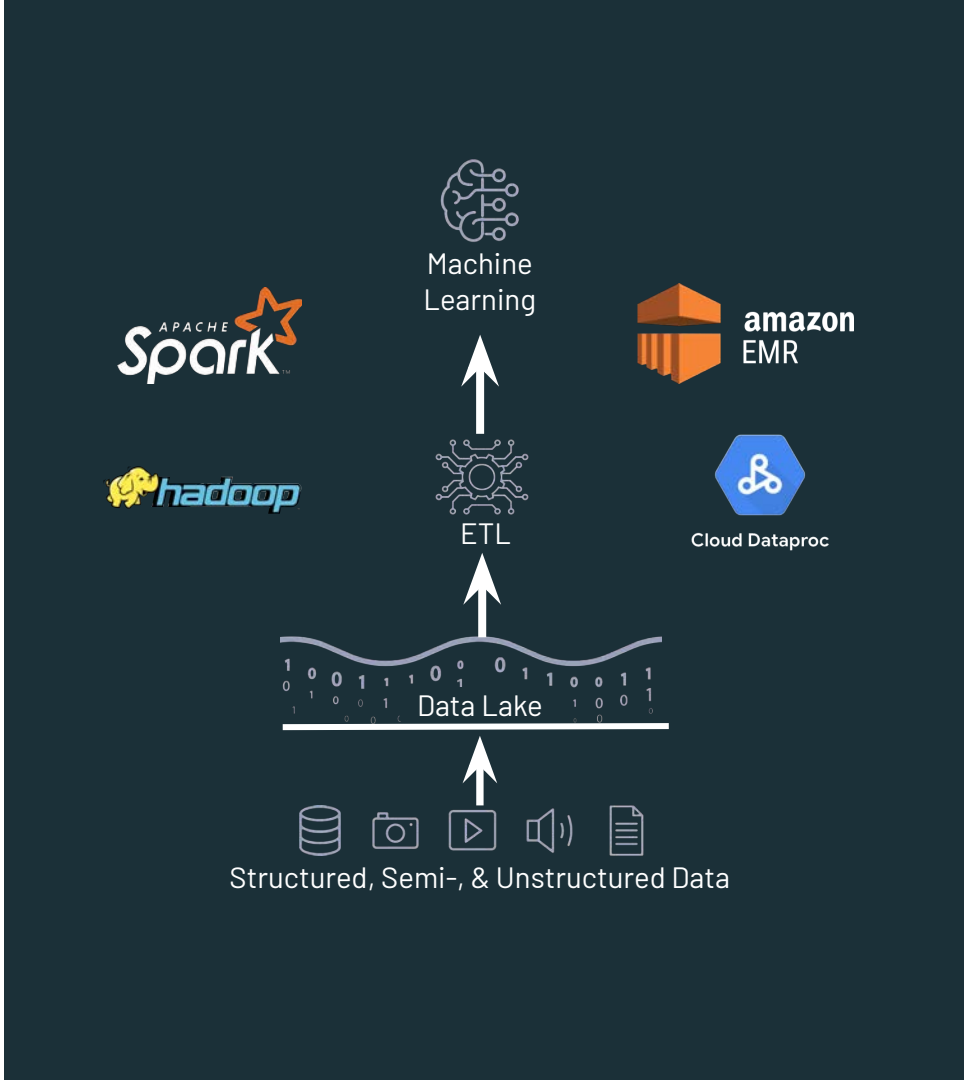

New Way Forward: Lakehouse

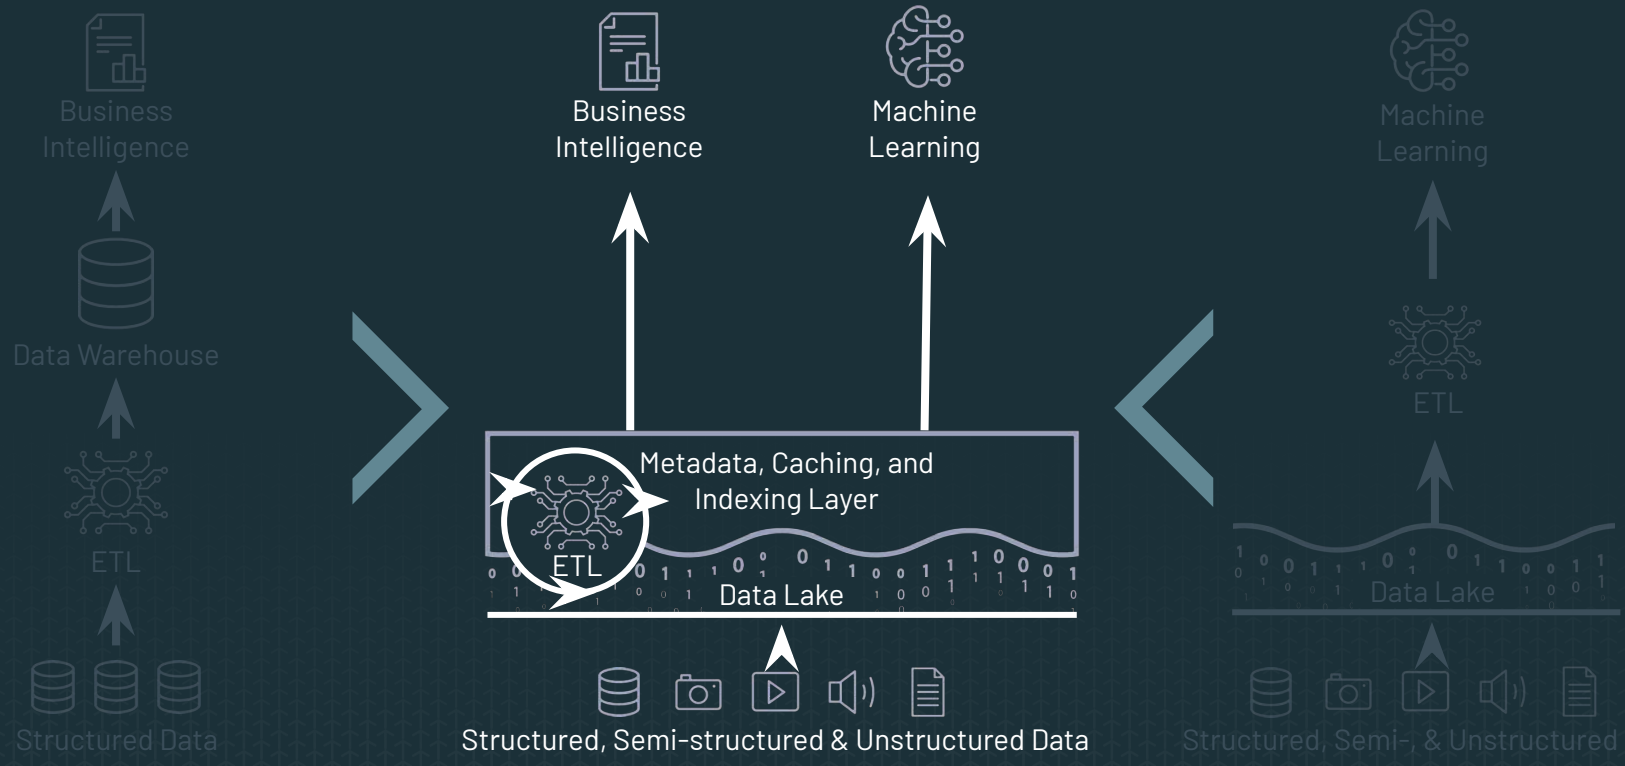

databricks

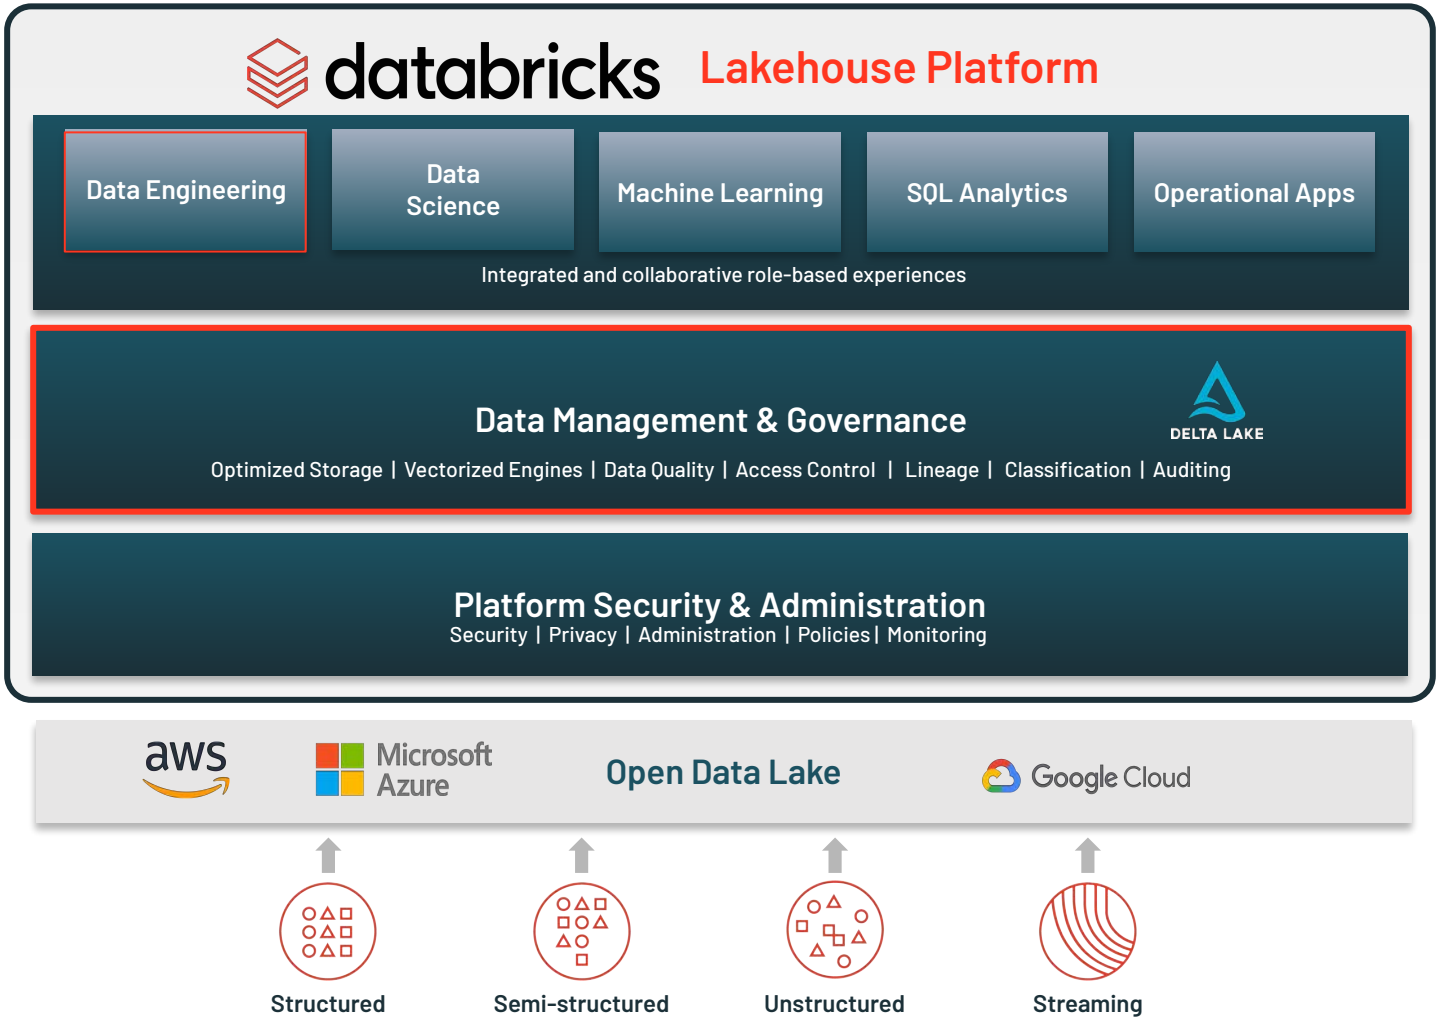
databricks
Lakehouse Platform
Data Engineering
Data Science
Machine Learning
SQL Analytics
Operational Apps
Integrated and collaborative role-based experiences
Data Management & Governance
DELTA LAKE
Optimized Storage | Vectorized Engines | Data Quality | Access Control | Lineage | Classification | Auditing
Platform Security & Administration
Security | Privacy | Administration | Policies | Monitoring
aws
Microsoft Azure
Open Data Lake
Google Cloud
Structured
Semi-structured
Unstructured
Streaming

Key differentiators for successful data engineering

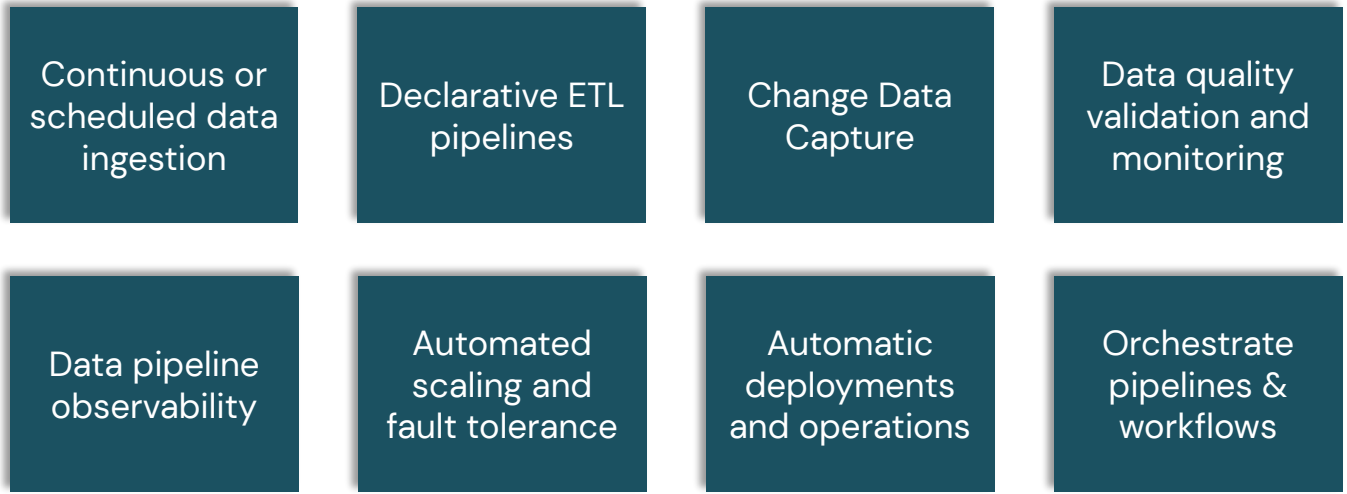

Declarative ETL pipelines with Delta Live Tables

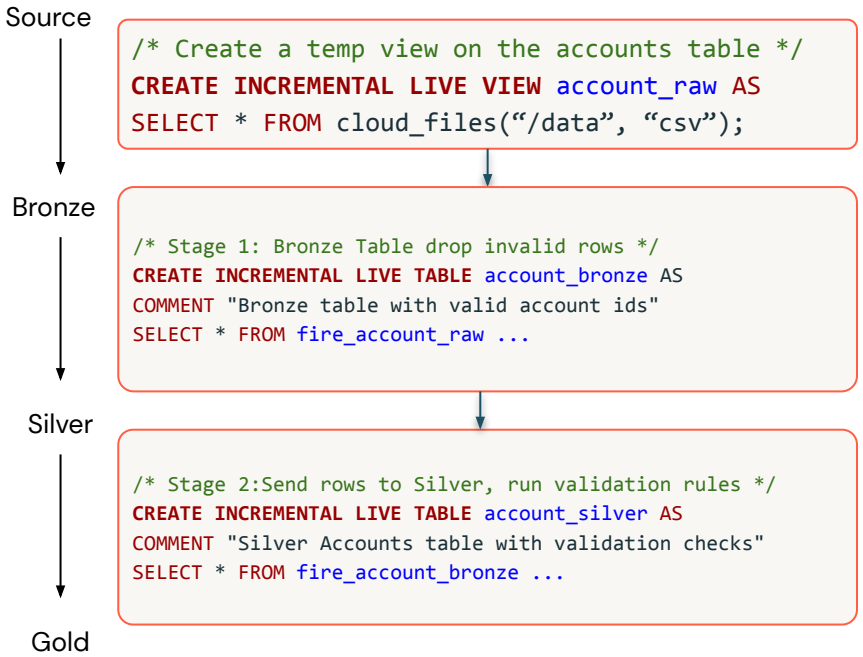

- Use intent-driven declarative development to abstract away the **"how"** and define **"what"** to solve
- Automatically create high-quality lineage and manage table dependencies across the data pipeline
- Automatically checks for errors, missing dependencies and syntax errors, and manage pipeline recovery

Continuous or scheduled data ingestion with Auto Loader

Simple SQL Syntax for Streaming Data Ingestion

Cmd 4

```
CREATE INCREMENTAL LIVE TABLE sales_orders_raw
COMMENT "The raw sales orders, ingested from /databricks-datasets."
TBLPROPERTIES ("quality" = "bronze")
AS
SELECT * FROM cloud_files
("/databricks-datasets/retail-org/sales_orders/",
"json", map("cloudFiles.inferColumnTypes", "true"));
```

- **Incrementally and efficiently** process new data files as they arrive in cloud storage
- Automatically **infer schema** of incoming files or superimpose what you know with **Schema Hints**
- Automatic **schema evolution**
- **Rescue data column** – never lose data again

			Coming Soon	Coming Soon
Schema Evolution	✅ JSON	✅ CSV	AVRO	PARQUET

Change data capture (CDC) with Delta Live Tables

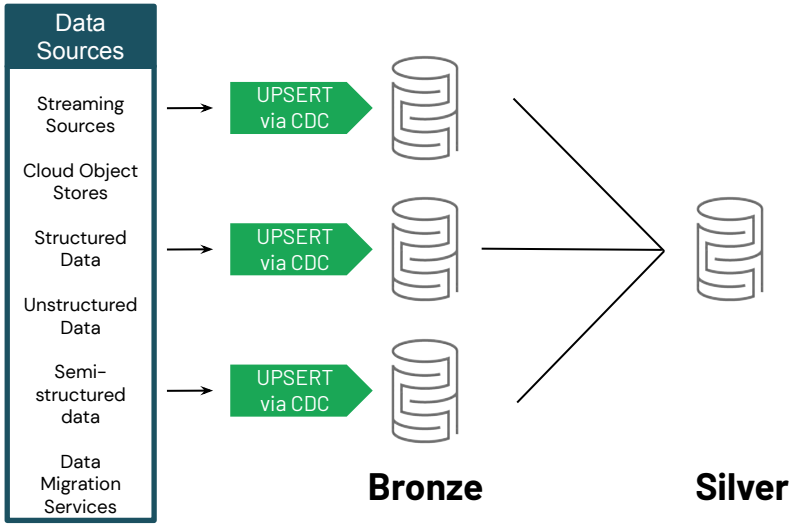

- Capture row-level changes from any data source supported by DBR, cloud storage, or DBFS
- Simpler architecture: build, simple incremental pipelines
- Handles out-of-order events
- Schema evolution
- Process change records (inserts, updates, deletes) incrementally using a simple, declarative “APPLY CHANGES INTO” SQL API

Automated scaling and fault tolerance with Delta Live Tables

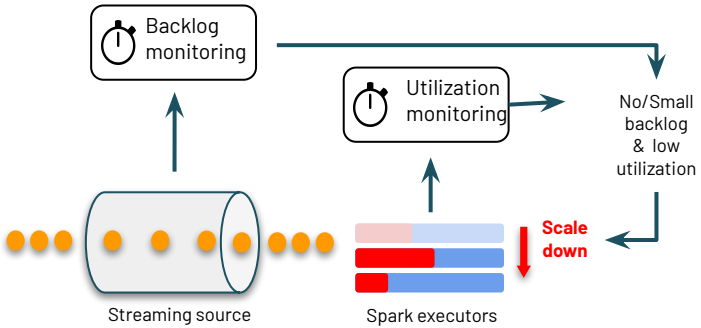

- Meet streaming SLOs with backlog-aware scaling decisions - Monitor both, **backlog metrics and cluster utilization** to scale up or down
- **Reduce down time** with automatic error handling and easy replay
- **Eliminate maintenance** with automatic optimizations of all Delta Live Tables
- Execute data pipeline workload on **automatically provisioned** elastic Apache Spark™-based compute clusters that parallelize jobs as well as minimize data movement

Automatic deployments and operations with Delta Lives Table

- Complete, parameterized and automated deployment for the continuous data delivery
- **Reuse ETL pipelines** across environments with config files and parameterization
- Orchestrates, tests, and monitor end-to-end the data pipeline

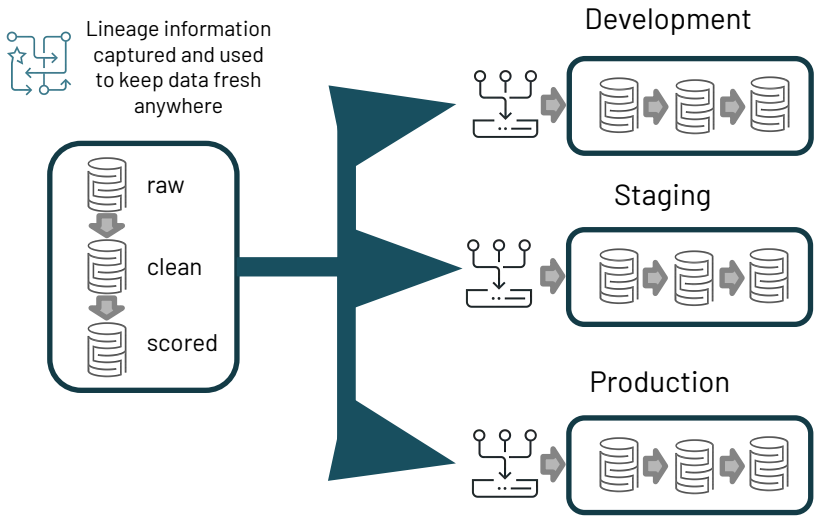

Workflow Management on Databricks

Simplify orchestration and management of multi-step workflows

Before

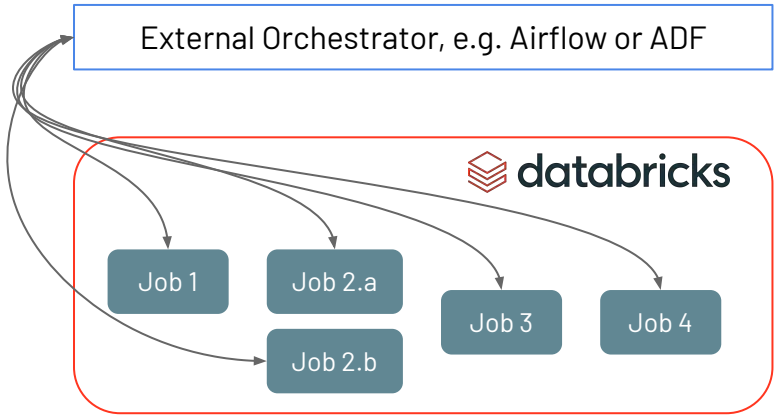

- Cost/complexity of maintaining external orchestrator
- Hard to monitor/debug

After

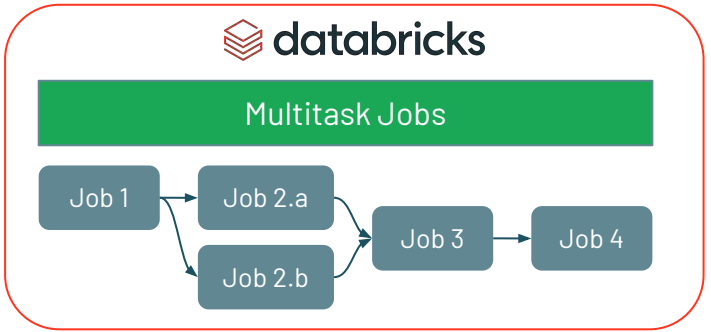

- Turnkey orchestration within Databricks
- Visibility into job dependencies, debugging, etc.
- Airflow and ADF integrations will continue to be supported

databricks

Demo

Additional Resources

- Getting Started with Delta Live Tables
- 5 Steps to Implementing Intelligent Data Pipelines With Delta Live Tables
- Product Page
- Documentation
- Spark's Structured Streaming
- Delta Lake
- Great Expectations